

AI Agents for Retail Operations

Part 0 of 3 · The venture strategy — where we play, how we win,
and how we pay for it

July 2026

1 · Strategy first

This is Part 0 of a three-part series, and it should be read first. We settled the strategy before we chose a product: a clear view of where a team like ours can win, and where it cannot. Every idea in the search was judged against that view. [Part 1](#) shows the research we did — 12 scans, 60+ candidate ideas, and the one product left standing at the end: **Buying Desk**, the purchase-order exception agent. [Part 2](#) lays out the plan to test that product idea before we build it. The three documents also have different lifespans: if the product fails its tests, Parts 1 and 2 get rewritten — this document should barely change.

Our starting point is an unusual advantage: ten years spent building and running the tech solutions of a real retail enterprise, including the buying desk. We know these problems from the inside, in the buying team's own words, before a prospect has told us anything.

And the strategy leads to this, in one sentence: a venture whose **AI agents do the manual work inside retail operations** — the chasing, checking, reconciling and data entry that mid-market retailers still do by hand out of shared inboxes and spreadsheets — bringing a human in only for the decisions that need real judgement. One rule is fixed from day one: **our products do the work; they never make the buying decisions.**

2 · Where we play — winnable ground, not exciting ground

We are a lean team, not a seed-funded startup. That single fact, stated up-front, drives the whole strategy.

The flashy AI categories — customer-service agents, e-commerce automation, sales tooling — are where startups that have raised \$20–100M out-spend each other for the same customers. Even if we built a superb product, it would not get far there: buyers hear from the companies with the biggest sales and marketing budgets, not from whoever built the best product, and a rival with ten times the engineers can copy any good feature within months. Competing there means entering a spending race with our own money against their investors' money.

So we deliberately chose different ground: **problems that hurt just as much but get none of the attention** — the unglamorous daily work inside retail operations that the big players overlook. The research in [Part 1](#) shows the one we picked (PO exceptions trapped in email and customised ERP estates) is real, recurring, costed by practitioners at six figures a year, and still unclaimed on the retail side — while the equivalent product for manufacturers has attracted over \$100M of funding, proving buyers pay for this job.

Two things make this ground winnable for a team like ours:

1. **The buyer is reached by knowledge, not by advertising.** A conservative, trust-gated retail buyer ignores cold SaaS marketing — which is exactly why the funded players haven't taken this ground yet. What opens the door is describing their problem better than they can themselves. We can.
2. **The economics fit a lean team.** No spending race, no free-tier war, no billboard budget. We will still have a strong go-to-market strategy — sales and marketing included — but on this ground, being heard takes focused outreach and trust, and that wins a customer at a fraction of what it costs in the crowded categories.

3 · How we win — the playbook

Sell the diagnosis before the cure. The entry product is a paid, read-only **Exception-Loss Audit**: two weeks, six weeks of the retailer's data, and a defensible £ number for what PO exceptions cost them — something most mid-market retailers cannot produce themselves. The audit is the demand test, the pilot and the first revenue in one, and it de-risks the technical build before we write production code.

Earn trust with a hard boundary. Everything is approval-gated: read-only first, human sign-off on every action, no writes to the ERP until the customer has watched the agent be right for weeks. No forecasting, no replenishment, no assortment decisions — ever. That keeps us out of the big incumbents' territory and inside the buyer's trust.

One agent first, a suite behind it. We sell one product — the PO-exception agent — solved to the core. The architecture is built for the agents that follow on the same desk, for the same buyer, through the same connector: cost-price-increase management, article/SKU setup, supplier onboarding. The suite is the roadmap, not the pitch.

Anchor deep, then widen. Smyths is the anchor design partner — a relationship we strengthen and protect throughout. New customers come from targeted outbound built on the domain story, then from references: in this market, one retailer vouching for us is worth more than any campaign.

Evidence gates everything. Every stage has pass and kill lines agreed before we start ([Part 2](#)). The same discipline that cut 60+ ideas down to one continues into how we build, sell and — as Section 5 sets out — how we finance.

Move inside our window. The funded players building this for manufacturers are realistically 12–18 months from crossing into retail. We don't need to beat their budgets; we need to be the ones already trusted by mid-market retail buyers when they arrive.

4 · What it costs, and what it earns

This is a **capital-light** business. The costs are a small team's time: validation costs conversations and one read-only data extract; the build is 2–3 engineers for roughly 6–9 months; there are no factories, no inventory, no

ad budgets. The technology is proven — document reading with large language models, deterministic comparison logic, email workflow, an approval cockpit — nothing speculative.

The revenue side, at the price point the research supports (~£35k per year per retailer, tested properly at [Gates 2 and 5](#)):

Milestone	Annual recurring revenue
First paying production customer	~£35k
10 customers	~£350k
30 customers	~£1M+

Alongside subscriptions, the paid audits generate revenue from the very first weeks — each one part-funds the next conversation. Once running, software of this kind carries high margins, so growth from around the first ten customers onward can plausibly pay for itself.

Two honest caveats. First, these are illustrations, not forecasts — the price point and the conversion rate are exactly what [Gates 2 and 5](#) exist to test. Second, capital-light does not mean free: the venture consumes real salaries for 12–18 months before subscription revenue covers them, and that carrying cost is the investment being made.

5 · Funding — when outside money makes sense, and when it doesn't

Start with a reality check on how “raisable” seed money actually is in 2026. The headlines say AI money is everywhere; the data says it is concentrated. Roughly four of every five venture dollars now go to AI — but overwhelmingly to US companies and to founders coming out of the big AI labs, with a handful of giants absorbing most of the total. For everyone else, seed has quietly become what Series A used to be: investors expect a working product, paying customers and typically \$300–500k of annual revenue before they write a cheque, and a typical UK software seed is £750k–£1.5m in exchange for 20–25% of the company. The practical conclusion cuts through the whole debate: **the evidence a team like ours would need**

to raise a seed round is the same evidence our gates produce. Whether or not we ever want outside money, the next step is identical — which is why the plan doesn't depend on answering the funding question today.

With that context, the default plan is to build this without outside investment. That is a strategy, not a limitation:

- Raising now would mean selling a slice of the company at its cheapest-ever price, before the evidence exists that makes it valuable.
- Investor money arrives with a clock attached. It pushes ventures toward exactly the land-grab behaviour we examined and rejected in Section 2 — spend fast, chase the crowded categories, buy growth.
- Our market cannot be bought faster. Trust-gated retail buyers move at the speed of references and results, not advertising. Extra millions would not shorten the sales cycle; they would just burn while we wait.
- Fundraising itself costs months of founder attention — the exact attention the validation programme needs.

“Couldn't we raise a seed round and go after one of the big categories instead?” It is the right question to ask, and the answer is no — for three reasons. First, the money wouldn't fix what blocks us there: a seed round is single-digit millions, while the players already in those categories have raised \$20–100M — we would arrive with the same problem and a fraction of the fuel, having sold a slice of the company for the ticket. Second, our real advantage doesn't travel: investors back teams with an unfair edge, and ours — a decade inside a retail buying desk, and the trust that comes with it — is worth a great deal on this ground and nothing in someone else's category, where we would be a team with no track record pitching against funded specialists. Third, even if funding were the goal, the sensible sequence would be unchanged: prove demand cheaply first, then raise on evidence at a far better price. Outside money can make a working engine go faster; it cannot buy us someone else's advantages.

When a seed round would start to make sense. Outside money becomes rational only when it scales a proven engine rather than funding a search. Concretely, all of the following, evidenced not hoped:

1. The five gates in [Part 2](#) have passed.
2. Two to three retailers are paying production subscriptions, with renewal signals.

3. The way we win a customer is repeatable and measured — audits convert to subscriptions at the rate Gate 5 demands.
4. The demand pipeline is larger than the current team can serve — money would buy engineers, implementation capacity and focused sales to serve customers we already know how to win.
5. The competitive window matters: if the funded manufacturing players start crossing into retail, speed becomes decisive, and outside capital is the fastest way to take the ground first.

A later, larger round (Series A) would only enter the conversation with a further layer of proof: customers renewing year on year, the second agent of the suite selling to existing customers, and unit economics that make “more capital in, predictably more revenue out” a calculation rather than a hope.

And we may choose never to raise. A capital-light, high-margin product that grows on references can fund its own growth. It is also worth knowing — as context, not as the plan — that acquirers exist in this category: ERP vendors have been buying exactly these companies (Kavida by QAD, Traverse by SPS Commerce). Every option stays open precisely because we will not be forced to raise.

Who decides, and when. The funding question is formally revisited at the 13-week gate review (the end of the validation programme in [Part 2](#)) and at each quarterly review after that — answered from the gate evidence on the table, never from enthusiasm or fear.

6 • The standing rules — what we will not do

1. **No building before the gates say build.** The validation programme comes first; its kill lines are real.
2. **No spending races.** If a category can only be won by outspending funded rivals, we don’t enter it — the kill-patterns in [Part 1](#) stand as permanent filters.
3. **No agent that makes buying decisions.** The trust boundary is the product’s identity.
4. **No raising money to search for product-market fit.** Outside capital scales proof; it never substitutes for it.
5. **Nothing that risks the Smyths relationship.** It is the venture’s foundation and is protected above any single deal or deadline.

7 · What happens next

The horizons below are a sequence, not calendar dates: Week 1 begins once we have committed to the programme and agreed who works on it.

Horizon	What happens	The decision at the end
Weeks 1–13	The validation programme (Part 2): Smyths engagement, outbound test, paid audits, read-only pilot	Build / re-plan / stop — and the first formal look at the funding question
Months 4–12	If gates pass: build audit pipeline → pilot-ready agent → approval-gated write-back; first paying production customer	Does the sales motion repeat beyond the anchor?
Months 12–18	Second and third production customers; agent #2 (cost-price-increase) scoped for existing customers	Scale from revenue, raise to accelerate, or hold — decided from evidence

The idea search that produced this plan is paused, not finished. If the gates kill the front-runner, the search reopens — with the map of 60+ evaluated ideas and eight kill-patterns as its head start, and this strategy document still standing: the ground we choose to fight on doesn't change with the first product.

Prepared July 2026. Part 0 of 3 — companion to [Part 1](#) (the research behind the product choice) and [Part 2](#) (the validation plan). Financial figures are illustrations at the researched price point, to be validated at Gates 2 and 5 before being treated as a plan.