

Buying Desk — The PO-Exception Agent for Retailers

Part 2 of 3 · The validation plan — five gates, Smyths as design partner, the 90-day programme

July 2026

[Part 0](#) set out the strategy; [Part 1](#) told the selection story: 12 scans, 60+ ideas, and why the retailer-side PO-exception agent was the one clean survivor. This part is the forward plan — what still has to be proven, how the Smyths engagement proves it, and the decisions we are asking of each other.

1 · What we still need to find out — the five gates

The idea search is paused with one front-runner standing — but selection is not yet conviction. Everything so far rests on public evidence: funding rounds, competitor products, practitioners' forum posts. Conviction has to come from inside a real buying team, and that starts with Smyths — their exception volumes, their hours spent chasing, their worst incidents, their £ number. That is exactly what the five gates below are designed to produce: each piece of remaining risk becomes a test with a pass line and a kill line agreed *before* we run it. In order:

Gate 1 — Distribution: can we reach buyers we don't already know? (6 weeks, hard deadline, runs first — before any code.) The goal: **three named design partners** — Smyths as the anchor, plus **two further retailers brought on board over the full 90 days**, needed as the second and third voices on the roadmap and as the first paid audits (Gate 2). The six-week test is the first step towards that goal, and it is deliberately about conversations, not signatures: can the audit offer book first conversations with buying or supply-chain directors at retailers where we have no history? The out-bound asset is what a decade of building and running a real buying desk's

tech solutions gave the team: we can describe a mid-market retailer's PO problems from the inside, in the buying team's own vocabulary, before they have said a word. Six weeks of focused outreach to a named target list is enough to learn whether that story opens doors. Every red-team review concluded the binding risk is distribution: a conservative, trust-gated buyer is the very reason this market is still open. This is also the cheapest test available — it costs conversations, not code. *Kill line: fewer than 4 first conversations booked from the six-week push → the channel isn't working; we stop and re-plan how we reach buyers before building anything.*

Gate 2 — Willingness to pay: does the paid audit sell? The entry product is a read-only, roughly two-week **Exception-Loss Audit**: we analyse six weeks of the retailer's PO, delivery and email data and hand back the annualised £ loss (as a defensible range), money-back if it finds less than its fee. It is the product wedge, the pilot and the demand test in one — because research showed most mid-market retailers *cannot* put a number on these losses themselves (in one study only 3 of 11 could). We sell the diagnosis before the cure. *Kill lines: fewer than 2 buyers across the first 10–20 audit conversations → stop. Audit finds less than its fee at a typical partner → the economics are wrong; refund and rethink.*

Gate 3 — How much of the problem is really ours? (*Measured inside every audit.*) Every exception the audit finds gets sorted into two piles. Some flow through the big electronic ordering networks (EDI providers such as SPS) — sooner or later the network companies will catch those themselves, cheaply, as a bundled feature. The rest live in emailed PDFs and the retailer's customised SAP estate, where no network can see them. Only the second pile is truly ours. This gate measures, at each customer, whether that pile is big enough to build a business on. *Kill line: more than 60% of the exception value sits in the networks' pile → our slice is too thin; stop before building.*

Gate 4 — Product truth: the two-week read-only pilot. (*No writes to the ERP; pass/kill lines agreed upfront.*) On the anchor's live data (top ~50 suppliers; PO, confirmation, delivery, goods-receipt and buying-mailbox feeds, read-only):

Metric	Pass line	If it misses
Share of real exceptions resolvable with no human touch	$\geq 40\text{--}50\%$	we don't ask them to subscribe
Reading accuracy on their suppliers' actual document formats	$\geq 95\%$	unread documents route to a human queue
Supplier replies matched to the right exception automatically	$\geq 70\%$	unmatched replies surface for a human
£ of loss found vs the audit fee	greater than the fee	fee refunded

These numbers are already printed inside the prototype as commitments — the dashboard ships empty and fills with the customer's measured numbers, never with our claims.

Gate 5 — Conversion: does the audit turn into a subscription? (*90-day window.*) The business model is audit → agent subscription (target ~£35k per customer per year; ~115 retailers ≈ £4–5M annual recurring revenue at maturity). The real risk here is that buyers love the diagnosis and stall on delegating the work. *Kill line: audit-to-subscription conversion below 30% within 90 days → re-plan pricing and packaging before scaling anything.*

Why gates rather than milestones: each one is built so that failure is cheap and informative. The most expensive mistake available to us is building for six months on an unproven channel. The entire programme above costs some weeks of conversations plus one read-only data extract.

2 · Smyths as design partner — the full question set

Smyths is the ideal anchor: mid-market scale, a large import/vendor base with a long non-EDI tail, seasonal promotion-driven buying, a customised SAP estate — and a buying team that already works every day in a cockpit

we built, with an approved write-path into SAP. The engagement has four workstreams. Each comes with the exact questions to ask and the validation it must produce.

2.1 · Discovery with the buying team — the pain, in their numbers

Before any interview, we mine our own side: after a decade of building the buying desk's tech solutions, we (the ZakApps team) already know which problems those tools solve and which are still worked around by hand. That catalogue becomes the interview guide here — and, later, the outbound story for the next retailers.

Then two or three structured interviews with buying operations and a Head Buyer. Open with the practitioner quotes from Part 1 (“does this happen here?”), not with a pitch.

- **Volume.** How many POs go out per year, to how many active suppliers? What share of suppliers confirm via EDI, versus email/PDF, versus portals, versus not at all?
- **Exception mix.** Per 100 POs, how often: no confirmation within the agreed window · price different from contract · partial quantity · date slip · short delivery at goods receipt? (Take their estimates first — the audit then measures reality, and the gap between the two is itself a selling point.)
- **The process today.** Where do confirmations land (a shared inbox?), who chases, tracked in what (spreadsheets?), how many hours per buyer per week? What do supplier contracts say (confirm within 48 hours? penalty clauses?) — and are penalties ever actually charged?
- **The money.** Can anyone put a £ figure on last year's overpaid invoices from confirmations above contracted cost, unclaimed short-delivery penalties, or promotion lines that missed launch because a PO silently went unconfirmed? (*Expected answer: no — that expectation is the audit's pitch. If they can answer, even better: it sizes the deal.*)
- **Worst incidents.** The last exception that really hurt — a stockout at a promotion launch, a big overpayment found late. Who noticed, when, and what did it cost?

- **Seasonality.** How does the exception load change August–November (the Christmas build)? Which suppliers and lines are promotion-critical? Toy retail gives us the sharpest urgency window: a pilot that proves itself before peak season sells itself.
- **The trust boundary.** What would they let an agent send to a supplier without reading it first — and what never? Which suppliers must always go to a human? (This calibrates the product’s autonomy controls.)

2.2 · Data and wedge-size validation — the audit itself

- **The access ask:** read-only extracts of six weeks (ideally six months) of POs, order confirmations, advance shipping notices, goods receipts, and price files / contract amendments — plus access to (or an export of) the buying-desk mailbox for the same window.
- **Measure the mix (Gate 3):** sort every exception found into Gate 3’s two piles — the networks’ and ours. Smyths’ number is the first real data point on whether our slice is big enough.
- **Data quality check (a known audit-breaker):** are contracted costs and amendments actually maintained in the systems, or do price files go stale? If stale, the audit reports ranges and says so — but we must know how bad it is *before* promising precision.
- **Parsing reality:** collect the actual confirmation formats of the top 50 suppliers (PDF layouts, email styles, portals). The $\geq 95\%$ reading-accuracy pass line gets tested against these, not against clean samples.
- **A baseline for the ROI story:** current manual touches per exception and time-to-resolve, so pilot improvements are measured against Smyths’ own pre-agent baseline instead of a vendor benchmark.

2.3 • Technical validation — the SAP estate and the integration path

- **Map the estate:** confirm exactly which SAP components are in play and which approved read paths exist for PO / confirmation / shipping / goods-receipt / price data — whether that is OData, BAPI, IDoc or extracts.
- **Reuse of the existing cockpit write-path:** the buying cockpit already writes to SAP through an approved route. Validate that the agent can — later, and always approval-gated — reuse that path for PO corrections. This is the integration advantage, so it must be confirmed, not assumed. Clean-core: no modifications to the SAP core.
- **Mailbox integration:** what will IT permit on the shared buying inbox — live read access, forwarding rules, or a batch export? (Read-only is enough for the audit and the pilot.)
- **The systems-integrator question:** Smyths' long-standing SAP partner is both a potential ally and a potential gatekeeper. Agree early how the read-only extract is provisioned and who signs it off, so the pilot doesn't stall in a queue we don't control.
- **Security and data handling:** what does Smyths require before supplier and pricing data leaves the estate (data-processing agreement, residency, retention)? Single-tenant data handling is already the product's design commitment — get their requirements in writing now.

2.4 • Commercial validation

- **The audit ask:** will Smyths pay for the Exception-Loss Audit — or, as the anchor, take it free in exchange for **reference rights** (a named case study, and taking reference calls from prospective customers once the pilot has proven itself)? Either answer is fine, but we need to know which: the next two partners must be asked to pay (that is Gate 2, and a free Smyths audit doesn't count towards it).
- **The economic buyer:** who signs — the Buying Director or Finance? Which budget line does “stop losing money on PO exceptions” come from? What approval threshold keeps this a one-signature purchase? (The whole easy-to-buy thesis depends on staying under it.)

- **The success definition:** what number, on whose dashboard, makes Smyths renew and act as a reference? Agree it before the pilot starts.
- **Pricing sanity check:** test the ~£35k/year price against the £ figure the audit finds. Too cheap signals unseriousness; too expensive breaks the one-signature sale.

3 · The 90-day programme

Week 1 begins once staffing and Smyths' backing are agreed — the programme is defined by its sequence, not by a calendar date.

Weeks	Workstream	Output / decision
1–2	Getting Smyths on board: buying-team interviews (§2.1); data access and data-processing agreement (§2.2/2.3); design-partner agreement	Signed anchor engagement; interview evidence pack
1–6	Gate 1 in parallel: the outbound push — first conversations with target retailers	Go / re-plan on the channel (fewer than 4 conversations = stop)
3–6	Exception-Loss Audit at Smyths on six weeks of data; Gate 3 classification inside it	The £ range + the wedge-size number; audit method hardened for the next retailers
6–10	Two-week read-only live pilot at Smyths (Gate 4 metrics); audits sold to the first outbound prospects (Gate 2)	Measured pass-line results; first paid-audit signals
10–13	Decision gate: review all five gates together	Build / re-plan / stop — with evidence, not enthusiasm

If the gates pass, the build order is already decided — and it is deliberately the reverse of the demo’s visual order: the audit pipeline first (smallest build, it *is* the paid product, and it de-risks the SAP connector), then the pilot-ready agent (detect → propose → approve → send → track), then hardening and the approval-gated write-back. Realistic path to the first paying production customer: roughly 6–9 months with a 2–3-engineer team. Nothing in the product needs speculative technology — it is document reading with LLMs, deterministic comparison logic, email workflow, and an approval cockpit. All of it exists today.

4 • What we need from each other

1. **Smyths' backing for the engagement** — an introduction at Buying-Director level, framed as “let’s measure what PO exceptions cost us”, plus the IT sign-off path for a read-only extract.
2. **The team’s decade of problem knowledge, made explicit.**
Structured sessions with our own people (the ZakApps team) who built and ran the buying desk’s tech solutions: which problems the tools solve, which are still worked around by hand, and what that says about every other mid-market retailer still running on spreadsheets and a shared inbox. This catalogue becomes the interview guide, the audit checklist and the outbound pitch — it is the single most decisive input the partnership can provide.
3. **A committed outbound push.** A named list of mid-market retailers (UK/Ireland first; 100+ active suppliers, heavy email/PDF confirmation tail, SAP or comparable ERP estate, a buying-ops team that chases manually), approached with the audit offer and the domain story above. Gate 1’s six weeks test whether the story books conversations; bringing two further design partners on board is the 90-day goal.
4. **Agreement to the kill lines.** The discipline that narrowed sixty ideas down to one only keeps its value if we hold the same standard now. If the outreach doesn’t book conversations, if the audits don’t sell, if the wedge measures thin, or if conversion stalls — we say so and re-plan. The gates are cheap; conviction without them is not.

Bottom line for this discussion: we have paused the idea search deliberately — not because this idea is a guaranteed winner, but because after twelve scans it is the strongest candidate on the board, and the next effort is better spent testing it than scanning further. If the gates kill it, the search reopens, with the map of dead ends already drawn. What we are asking each other to commit to now is not “build the product”. It is a 13-week, evidence-gated validation programme whose main cost is a focused outbound push and one read-only data extract — and whose output is either a validated, defensible product thesis with an anchor customer, or a cheap, well-documented “no” that saves us a year.

Prepared July 2026. Part 2 of 3 — companion to [Part 0](#) (the strategy) and [Part 1](#) (the research behind the product choice). Companion materials: the Buying Desk clickable prototype · the plain-language product write-up · the PRD and suite strategy (on request).